



Chapter 3.

Processes 1.



**Operating System
Concepts (10th Ed.)**





3.1 Process Concept

- A *process* is a program in execution.
 - A process is the unit of work in an operating system.
 - A process will need certain resources to accomplish its task.
 - CPU time,
 - memory,
 - files,
 - and I/O devices.



3.1 Process Concept

- The memory layout of a process is divided into multiple sections:
 - Text section:
 - the executable code
 - Data section:
 - global variables
 - Heap section:
 - memory that is dynamically allocated during program run time
 - Stack section:
 - temporary data storage when invoking functions
 - such as function parameters, return addresses, and local variables



3.1 Process Concept




Figure 3.1 *Layout of a process in memory.*



3.1 Process Concept

```
#include <stdio.h>
#include <stdlib.h>

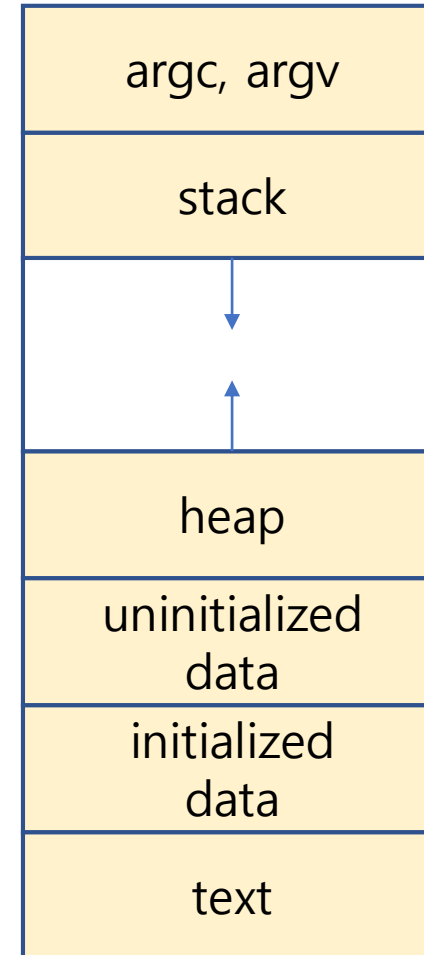
int x;
int y = 15;

int main(int argc, char *argv[])
{
    int *values;
    int i;

    values = (int *)malloc(sizeof(int)*5);

    for (i = 0; i < 5; i++)
        values[i] = i;

    return 0;
}
```






3.1 Process Concept

```
$ gcc 3.1_memory_layout.c
```

```
$ size ./a.out
```

text	data	bss	dec	hex	filename
1603	604	12	2219	8ab	./a.out





3.1 Process Concept

- As a process executes, it changes its **state**.
 - **New**: the process is being created.
 - **Running**: Instructions are being executed.
 - **Waiting**: the process is waiting for some event to occur.
 - such as an I/O completion or reception of a signal.
 - **Ready**: the process is waiting to be assigned to a processor.
 - **Terminated**: the process has finished execution.



3.1 Process Concept

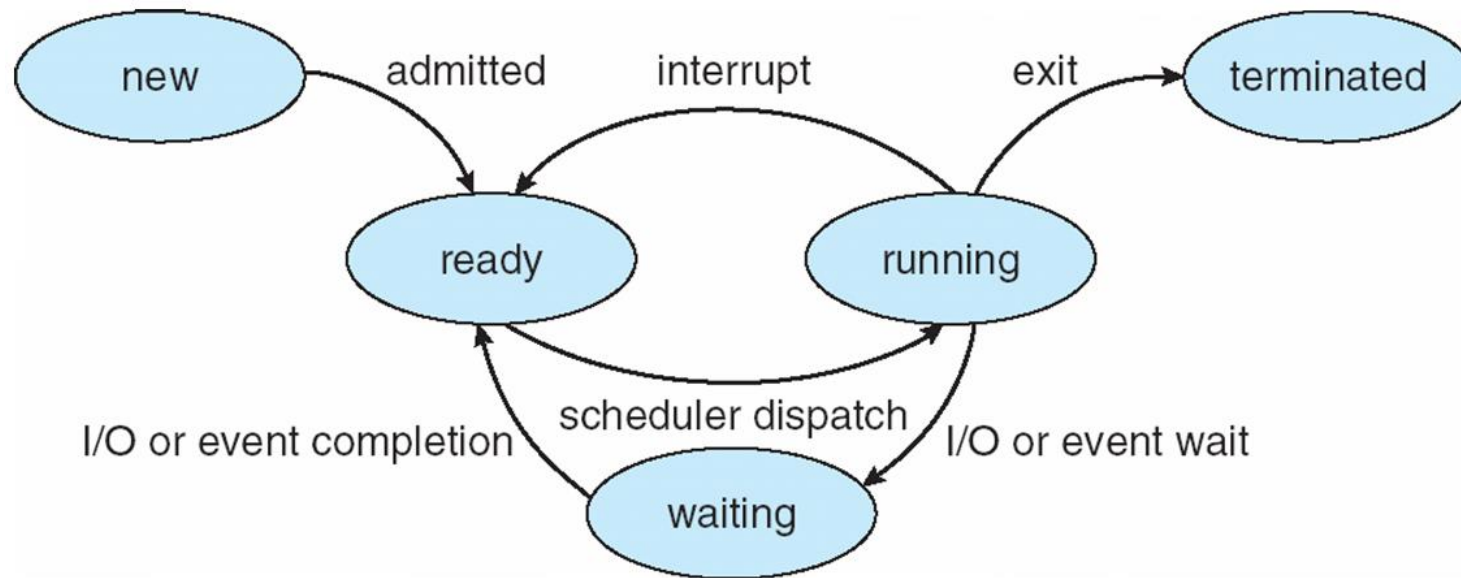


Figure 3.2 *Diagram of process state.*



3.1 Process Concept

- **PCB (Process Control Block)** or TCB (Task Control Block)
 - Each process is represented in the operating system by the PCB.
- A PCB contains many pieces of information associated with a specific process:
 - **Process state**
 - **Program counter**
 - CPU registers
 - CPU-scheduling information
 - Memory-management information
 - Accounting information
 - I/O status information



3.1 Process Concept

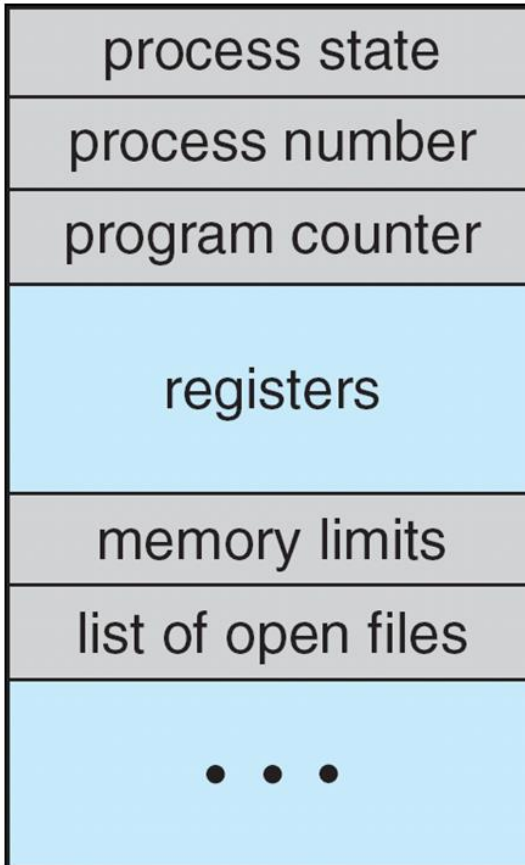


Figure 3.3 *Process control block (PCB).*



3.1 Process Concept

- A process is
 - a program that performs a ***single thread of execution***.
 - The single thread of control allows the process to perform
 - only one task at a time.
 - Modern operating systems have extended the process concept
 - to allow a process to have multiple threads of execution
 - and thus to perform more than one task at a time.
- A **thread** is a *lightweight* process.
 - Chapter 4 explores multithreading in detail.



3.2 Process Scheduling

- The objective of **multiprogramming** is
 - to have some process running at all times
 - so as to maximize CPU utilization.
- The objective of **time sharing** is
 - to switch a CPU core among processes so frequently
 - that users can interact with each program while it is running.



3.2 Process Scheduling

■ Scheduling Queues:

- As processes enter the system, they are put into a **ready queue**,
 - where they are ready and waiting to execute on a CPU's core.
- Processes that are waiting for a certain event to occur
 - are placed in a **wait queue**.
- These queues are generally implemented
 - in the linked lists of PCBs.



3.2 Process Scheduling

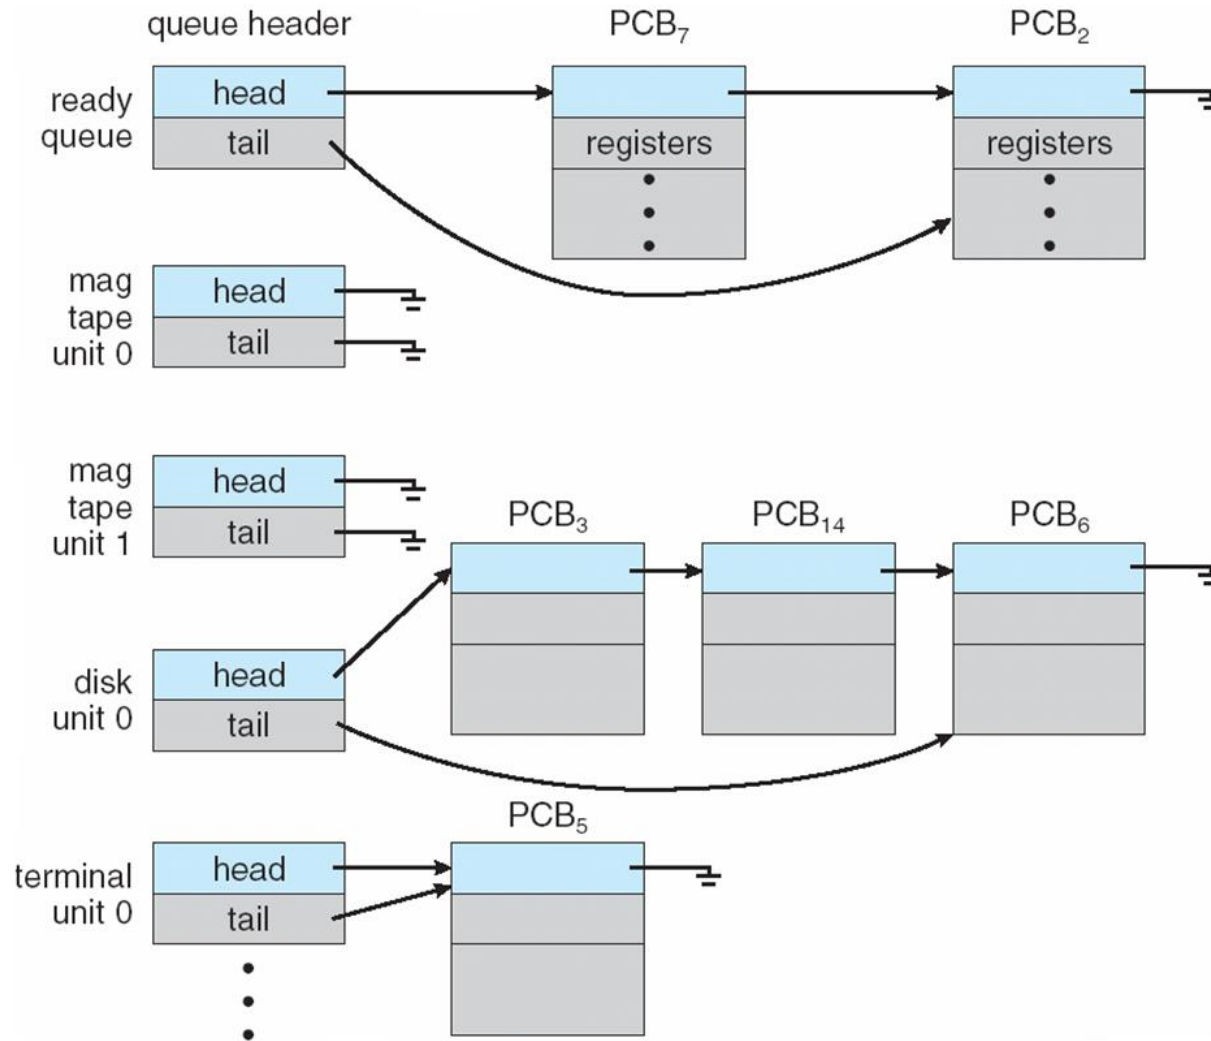


Figure 3.4 *The ready queue and wait queues.*



3.2 Process Scheduling

■ Queueing Diagram

- as a common representation of process scheduling.

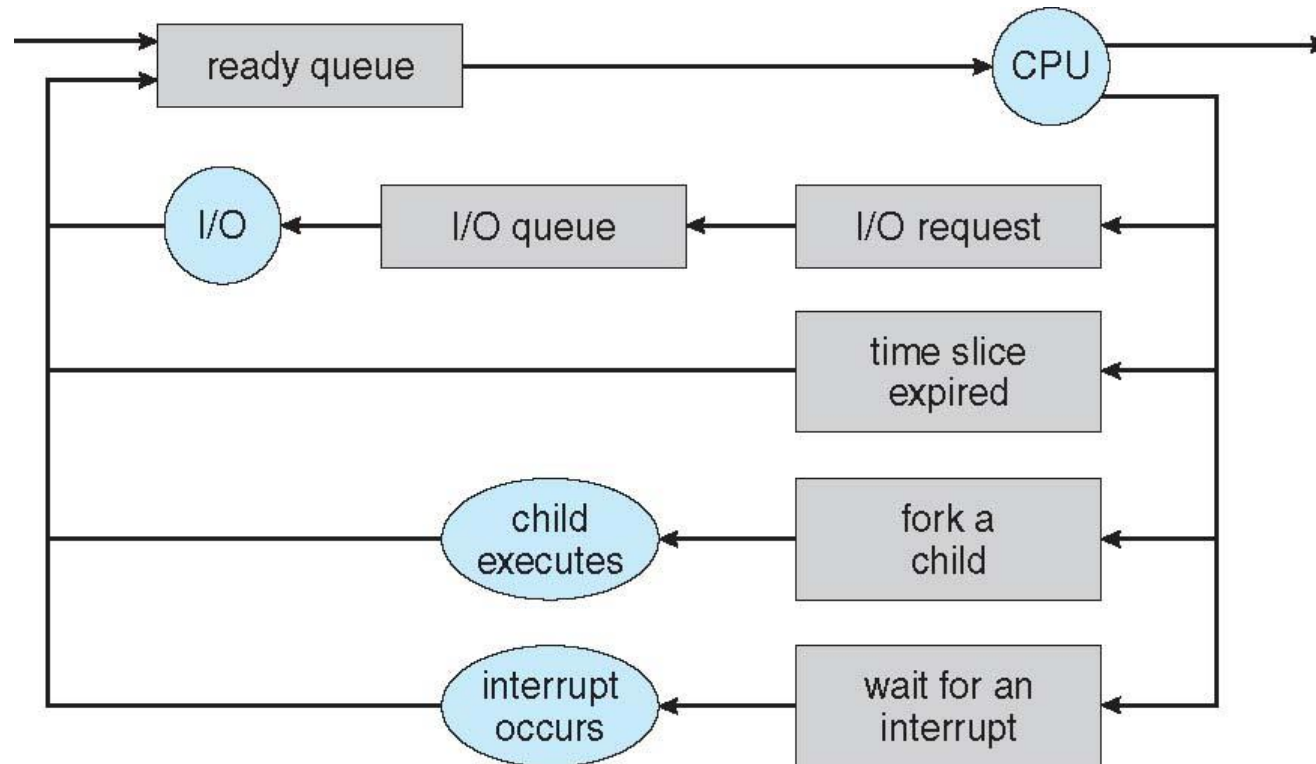


Figure 3.5 Queueing-diagram representation of process scheduling.



3.2 Process Scheduling

■ Context Switch

- The **context** of a process is represented in the PCB.
- When an interrupt occurs,
 - the system **saves** the current **context** of the running process,
 - so that, later, it can **restore** that **context** when it should be resumed.
- The **context switch** is a task that
 - switches the CPU core to another process.
 - performs a *state save* of the current process
 - and a *state restore* of a different process.



3.2 Process Scheduling

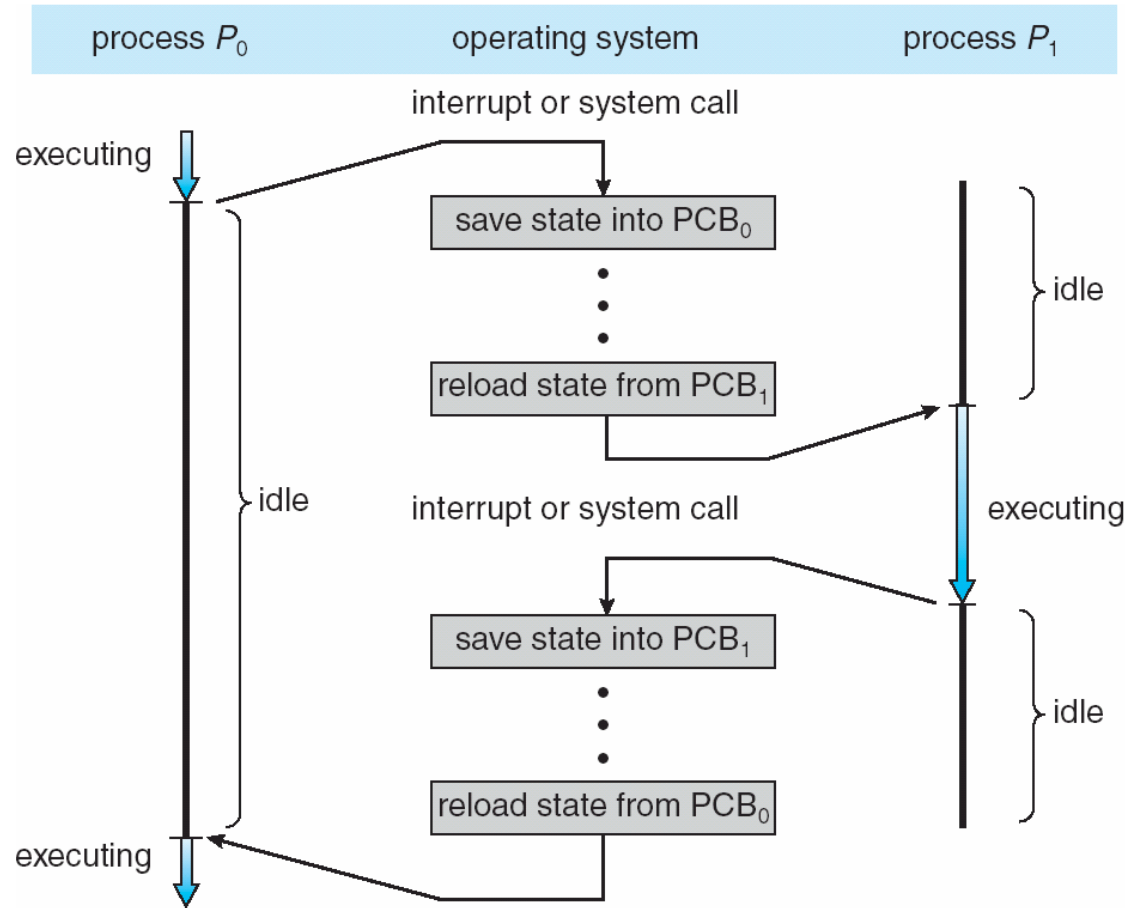


Figure 3.6 Diagram showing context switch from process to process.



3.3 Operations on Processes

- An operating system must provide a mechanism for
 - process creation,
 - and process termination.

- A process may create several new processes
 - the creating process: a *parent* process.
 - a newly created process: a *child* process.



3.3 Operations on Processes

■ A *tree* of processes

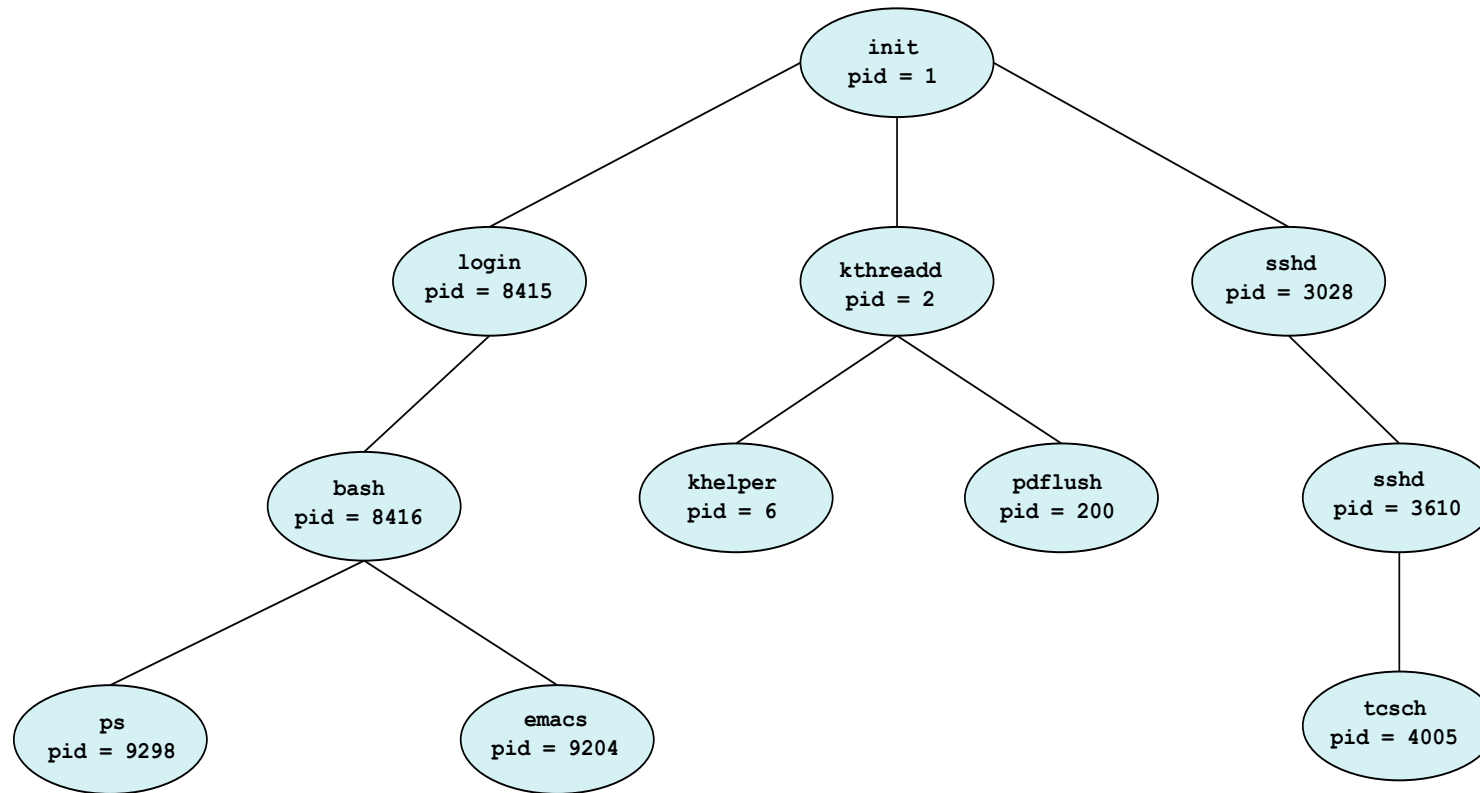


Figure 3.7 A tree of processes on a typical Linux system.



3.3 Operations on Processes

- Two possibilities for execution
 - The parent continues to ***execute concurrently*** with its children.
 - The parent ***waits*** until some or all of its children have terminated.

- Two possibilities of address-space
 - The child process is a ***duplicate*** of the parent process.
 - The child process has a ***new program*** loaded into it.



3.3 Operations on Processes

```
#include <stdio.h>      int main()
#include <unistd.h>    {
#include <wait.h>      pid_t pid;
                      // fork a child process
                      pid = fork();
                      if (pid < 0) { // error occurred
                          fprintf(stderr, "Fork Failed");
                          return 1;
                      }
                      else if (pid == 0) { // child process
                          execlp("/bin/ls", "ls", NULL);
                      }
                      else { // parent process
                          wait(NULL);
                          printf("Child Complete");
                      }
                      return 0;
                }
```

Figure 3.8 Creating a separate process using the UNIX fork() system call.



3.3 Operations on Processes

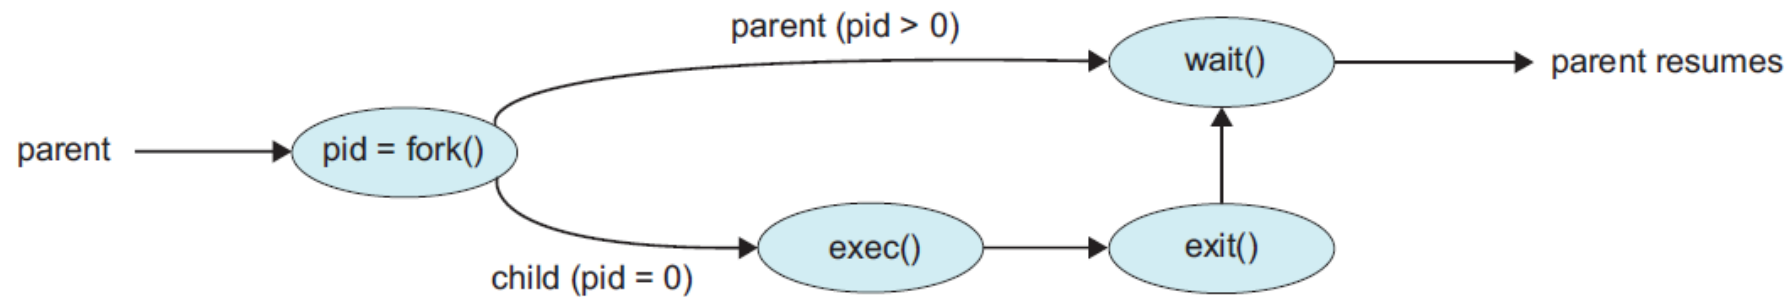


Figure 3.9 Process creation using the `fork()` system call.



3.3 Operations on Processes

- A process *terminates*
 - when it finishes executing its final statement
 - `exit()` system call: asks OS to delete it.
 - OS deallocates and reclaims all the resources:
 - allocated memories, open files, and I/O buffers, etc.



3.3 Operations on Processes

■ Zombie and Orphan

- *zombie* process: a process that has terminated,
 - but whose parent has not yet called wait().
- *orphan* process: a process that has a parent process
 - who did not invoke wait() and instead terminated.



3.3 Operations on Processes

- In UNIX-like O/S,
 - A new process is created by the **fork()** system call.
 - The *child* process consists of
 - a ***copy of the address space*** of the *parent* process.
 - Both processes continue execution
 - at the instruction after the fork() system call.
 - With one difference:
 - the return code for the fork() is **zero** for the child process, whereas
 - the **nonzero pid** of the child is returned to the parent process.



3.3 Operations on Processes

```
#include <stdio.h>
#include <unistd.h>

int main()
{
    pid_t pid;
    pid = fork();
    printf("Hello, Process!\n");
}
```



3.3 Operations on Processes

```
#include <stdio.h>
#include <unistd.h>

int main()
{
    pid_t pid;
    pid = fork();
    printf("Hello, Process! %d\n", pid);
}
```



3.3 Operations on Processes

- After a `fork()` system call,
 - the parent can ***continue its execution***; or
 - if it has nothing else to do while the child runs,
 - it can issue a **`wait()`** system call
 - to move itself off the ready queue until the termination of the child.



3.3 Operations on Processes

```
#include <stdio.h>
#include <unistd.h>
#include <wait.h>

int main()
{
    pid_t pid;
    pid = fork();
    if (pid > 0)
        wait(NULL);
    printf("Hello, Process! %d\n", pid);
}
```



3.3 Operations on Processes

■ Exercise 3.1 (p. 154)

```
int value = 5;

int main()
{
    pid_t pid;
    pid = fork();

    if (pid == 0) { // child process
        value += 15;
        return 0;
    }
    else if (pid > 0) { // parent process
        wait(NULL);
        printf("Parent: value = %d\n", value); // LINE A
    }
}
```

Figure 3.30 *What output will be at Line A?*



3.3 Operations on Processes

■ Exercise 3.2 (p. 154)

```
#include <stdio.h>
#include <unistd.h>
#include <wait.h>

/*
 * How many processes are created?
 */
int main()
{
    fork(); // fork a child process
    fork(); // fork another child process
    fork(); // and fork another

    return 0;
}
```

Figure 3.31 *How many processes are created?*



3.3 Operations on Processes

■ Exercise 3.11 (p. 905)

```
#include <stdio.h>
#include <unistd.h>

/*
 * How many processes are created?
 */
int main()
{
    int i;

    for (i = 0; i < 4; i++)
        fork();

    return 0;
}
```

Figure 3.32 *How many processes are created?*



3.3 Operations on Processes

■ Exercise 3.12 (p. 905)

```
int main()
{
    pid_t pid;
    pid = fork();

    if (pid == 0) { // child process
        execlp("/bin/ls", "ls", NULL);
        printf("LINE J\n");
    }
    else if (pid > 0) { // parent process
        wait(NULL);
        printf("Child Complete\n");
    }

    return 0;
}
```

Figure 3.33 *When will LINE J be reached?*



3.3 Operations on Processes

■ Exercise 3.13 (p. 905)

```
int main()
{
    pid_t pid, pid1;
    pid = fork();
    if (pid == 0) { // child process
        pid1 = getpid();
        printf("child: pid = %d\n", pid); // A
        printf("child: pid1 = %d\n", pid1); // B
    }
    else if (pid > 0) { // parent process
        pid1 = getpid();
        printf("child: pid = %d\n", pid); // C
        printf("child: pid1 = %d\n", pid1); // D
        wait(NULL);
    }
    return 0;
}
```

Figure 3.34 *What are the pid values?*



3.3 Operations on Processes

■ Exercise 3.16 (p. 905)

```
#define SIZE 5
int nums[SIZE] = {0, 1, 2, 3, 4};

int main()
{
    pid_t pid;
    int i;
    pid = fork();

    if (pid == 0) { // child process
        for (i = 0; i < SIZE; i++) {
            nums[i] *= i;
            printf("CHILD: %d \n", nums[i]); // LINE X
        }
    }
    else if (pid > 0) { // parent process
        wait(NULL);
        for (i = 0; i < SIZE; i++) {
            printf("PARENT: %d \n", nums[i]); // LINE Y
        }
    }

    return 0;
}
```

Figure 3.35 What output will be at Line X and Line Y?

Any Questions?

